
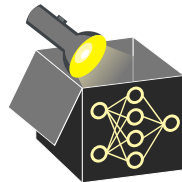


Interpretable Machine Learning

Explainable Boosting Machines (EBM)


Interpretable Models 1



Learning goals

- Understand link between GAM and EBM
- Learn univariate EBMs
 $\hat{=}$ GAM + boosting + shallow bagged trees
- Extend to GA2M: GAMs with selected pairwise interactions
- Detect interactions efficiently using FAST algorithm

RECAP: SPLIT SELECTION DECISION TREE



- **Impurity (Regression):** Variance of target Y in a node:

$$\text{Var}(Y) = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \bar{y})^2 = \frac{1}{n} \sum_{i=1}^n (y^{(i)})^2 - \bar{y}^2$$

- **Sum of squared errors (SSE) = residual sum of squares (RSS):**

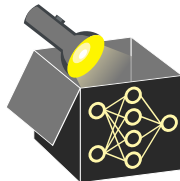
$$\text{RSS} = n \cdot \text{Var}(Y) = \sum_{i=1}^n (y^{(i)} - \bar{y})^2 = \dots = \sum_{i=1}^n (y^{(i)})^2 - \frac{1}{n} \left(\sum_{i=1}^n y^{(i)} \right)^2$$

Hence: $\text{RSS} = SS_n - \frac{S_n^2}{n}$ with $S_n = \sum_{i=1}^n y^{(i)}$, $SS_n = \sum_{i=1}^n (y^{(i)})^2$

- **Split criterion:**

- **Minimize post-split RSS:** $\text{RSS}_{\text{split}} = \text{RSS}_L + \text{RSS}_R$
- **Maximize reduction in RSS:** $\Delta \text{RSS} = \text{RSS}_{\text{parent}} - (\text{RSS}_L + \text{RSS}_R)$

NAIVE SPLIT SELECTION: EXPLICIT COMPUT.




- For a given feature X_j , sort the pairs $(x_j^{(i)}, y^{(i)})$ by increasing $x_j^{(i)}$.
- For each of the $n - 1$ potential split points at $s_k = \frac{1}{2}(x_j^{(k)} + x_j^{(k+1)})$:
 - Define partitions: $\mathcal{I}_L = \{i : x^{(i)} \leq s_k\}$, $\mathcal{I}_R = \{i : x^{(i)} > s_k\}$
 - Compute group means and counts after splitting at s_k :

$$\bar{y}_L = \frac{1}{N_L} \sum_{i \in \mathcal{I}_L} y^{(i)}, \quad \bar{y}_R = \frac{1}{N_R} \sum_{i \in \mathcal{I}_R} y^{(i)}, \quad \text{with } N_L = |\mathcal{I}_L|, \quad N_R = |\mathcal{I}_R|$$

- Compute RSS after splitting at s_k :

$$\text{RSS}_{\text{split}}(s_k) = \text{RSS}_L(s_k) + \text{RSS}_R(s_k) = \sum_{i \in \mathcal{I}_L} (y^{(i)} - \bar{y}_L)^2 + \sum_{i \in \mathcal{I}_R} (y^{(i)} - \bar{y}_R)^2$$

- Select split point s_k that minimizes $\text{RSS}_{\text{split}}(s_k)$
- **Compute cost:** $O(n^2)$ per feat. (recompute mean & RSS at each split)



$O(n^2)$ operations (recompute for each split s_j per feature)

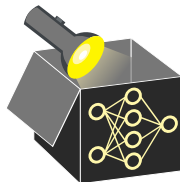
EFFICIENT SPLIT SELECTION

- **Setup:** For feature X_j , sort the data $(x_j^{(i)}, y^{(i)})_{i=1}^n$ by increasing $x_j^{(i)}$
- **Define group statistics (cumulative sums) after split at s_k :**


$$S_L = \sum_{i \in \mathcal{I}_L} y^{(i)}, \quad SS_L = \sum_{i \in \mathcal{I}_L} (y^{(i)})^2, \quad N_L = |\mathcal{I}_L|$$
$$S_R = S_n - S_L, \quad SS_R = SS_n - SS_L, \quad N_R = n - N_L$$

- **RSS for child nodes and parent node:**

$$\text{RSS}_L(s_k) = SS_L - \frac{S_L^2}{N_L}, \quad \text{RSS}_R(s_k) = SS_R - \frac{S_R^2}{N_R}, \quad \text{RSS}_{\text{parent}} = SS_n - \frac{S_n^2}{n}$$



EFFICIENT SPLIT SELECTION



- **Setup:** For feature X_j , sort the data $(x_j^{(i)}, y^{(i)})_{i=1}^n$ by increasing $x_j^{(i)}$
- **Define group statistics (cumulative sums) after split at s_k :**

$$S_L = \sum_{i \in \mathcal{I}_L} y^{(i)}, \quad SS_L = \sum_{i \in \mathcal{I}_L} (y^{(i)})^2, \quad N_L = |\mathcal{I}_L|$$
$$S_R = S_n - S_L, \quad SS_R = SS_n - SS_L, \quad N_R = n - N_L$$

- **RSS for child nodes and parent node:**

$$\text{RSS}_L(s_k) = SS_L - \frac{S_L^2}{N_L}, \quad \text{RSS}_R(s_k) = SS_R - \frac{S_R^2}{N_R}, \quad \text{RSS}_{\text{parent}} = SS_n - \frac{S_n^2}{n}$$


- **Reduction in RSS:**

$$\Delta \text{RSS}(s_k) = \text{RSS}_{\text{parent}} - (\text{RSS}_L + \text{RSS}_R) = \frac{S_L^2}{N_L} + \frac{S_R^2}{N_R} - \frac{S_n^2}{n}$$

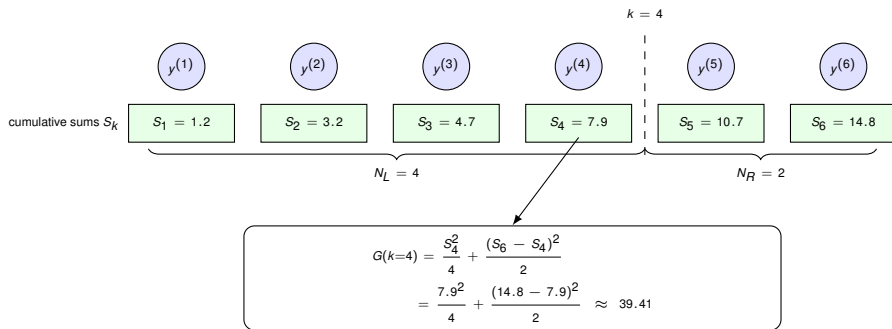
All squared-target terms SS_L, SS_R cancel. Only first-order sums needed.

- **Search:** Choose best split $s_k^* = \arg \max_{s_k} \Delta \text{RSS}(s_k)$
- **Complexity per feature:**
 $O(n \log n)$ (sorting) + $O(n)$ (cumulative sums and scan)

EFFICIENT SPLIT SELECTION - EXAMPLE



$$y^{(1)} = 1.2, y^{(2)} = 2.0, y^{(3)} = 1.5, y^{(4)} = 3.2, y^{(5)} = 2.8, y^{(6)} = 4.1 \quad (x_j^{(1)} \leq \dots \leq x_j^{(6)})$$



- $G(k)$ omits $-S_n^2/n$ (identical for all splits \Rightarrow does not affect arg max).
- Only cumulative sums S_k are required, no SS_k is stored or updated.
- $\mathcal{O}(1)$ per split $\Rightarrow \mathcal{O}(n)$ per feature.

EXPLAINABLE BOOSTING MACHINES (EBM)


Recall GAM:

$$g(\mathbb{E}[y \mid \mathbf{x}]) = \theta_0 + f_1(x_1) + f_2(x_2) + \dots + f_p(x_p),$$

- One shape function f_j per feature x_j
~> **Feature-level interpretability**
- Captures non-linear univariate effects
~> **Better performance / more flexible than GLMs**

EBM idea: GAMs train with **gradient boosting** over **shallow bagged trees**

- **GAMs** - feature-wise interpretability via separate shape functions $f_j(x_j)$
~> Potentially include pairwise interactions manually
- **Gradient Boosting** - incrementally fits residuals to improve predictive performance while retaining additivity
- **Shallow Bagged Trees** - low-depth trees (2–4 leaves) reduce variance and create interpretable shape functions




EBM - TWO-STAGE MODEL CONSTRUCTION

1 Stage 1: Fit Main Effects (Univariate Terms) ▶ Lou 2012

- Train EBM using only feature-wise shape functions $f_j(x_j)$
- Freeze the univariate model after convergence

2 Stage 2: Add Selected Pairwise Interactions ▶ Lou 2013

- Apply **FAST** to rank all $O(p^2)$ feat pairs by potential reduction in RSS
- Select top K pairwise interactions and store them in \mathcal{K}
- Use boosting to fit pairwise interaction terms $f_{ij}(x_i, x_j)$ on residuals
- Final model: $\hat{f}() = \sum_{j=1}^p f_j(x_j) + \sum_{(i,j) \in \mathcal{K}} f_{ij}(x_i, x_j)$



UNIVARIATE EBM - INITIALIZATION

- Set all shape functions to zero:


$$f_j^{[0]}(x_j) = 0 \quad \text{for all } j = 1, \dots, p$$

- Compute initial model prediction:

$$\hat{y}^{[0]} = \sum_{j=1}^p f_j^{[0]}(x_j) = 0$$

- Compute initial pseudo-residuals (e.g., for squared loss):

$$\tilde{r}^{[0]} = -\frac{\partial L}{\partial \hat{y}} = y - \hat{y}^{[0]} = y$$



UNIVARIATE EBM FIRST FEATURE UPDATE

Iteration feat_1 feat_2 feat_3 ... feat_p

1



$\xrightarrow{\text{res}_1}$

- Fit shallow bagged tree $T_1^{[1]}$ (2–4 leaves) to training data $\left\{ (x_1, \tilde{r}^{[0]})^{(i)} \right\}_{i=1}^n$
 \rightsquigarrow Use only feature x_1 as input and $\tilde{r}^{[0]}$ as target
- Update first shape function with learning rate η :


$$f_1^{[1]}(x_1) = f_1^{[0]}(x_1) + \eta \cdot T_1^{[1]}(x_1)$$

- Update prediction:



$$\hat{y}^{[1]} = \sum_{j=1}^p f_j^{[1]}(x_j)$$

- Recompute pseudo-residuals:

$$\tilde{r}^{[1]} = -\frac{\partial L}{\partial \hat{y}} = y - \hat{y}^{[1]}$$



UNIVARIATE EBM CYCLE THROUGH FEATURES





- 1st boosting iteration:

Cycle through each feature $j = 2, \dots, p$:


- Fit shallow bagged tree $T_j^{[1]}$ using feature x_j and previous residual $\tilde{r}^{[j-1]}$
 - Update f_j : $f_j^{[1]}(x_j) = f_j^{[0]}(x_j) + \eta \cdot T_j^{[1]}(x_j)$
 - Recompute \hat{y} and residuals: $\tilde{r}^{[j]} = y - \hat{y}^{[j]}$
- After one full pass over features, we complete one boosting iteration

UNIVARIATE EBM ITERATE BOOSTING PROCESS



- Repeat feature-wise updates for M boosting iterations (e.g., $M = 10000$)
- In each boosting iteration:
 - Cycle over all features $j = 1, \dots, p$ individually
 - Update only one f_j at a time using residuals from previous state
- Use small learning rate η to ensure smooth updates and order-invariance

UNIVARIATE EBM - PREDICTION & INTERPRETABILITY




- Final model consists of M shallow trees per feature:

$$\text{EBM Model} = \sum_{j=1}^p \sum_{m=1}^M \eta \cdot T_j^{[m]}(x_j)$$

- For each feature x_j , combine its M trees into a shape function:

$$\hat{f}_j(x_j) = \sum_{m=1}^M \eta \cdot T_j^{[m]}(x_j)$$

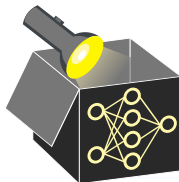
- Plot $\hat{f}_j(x_j)$ vs. $x_j \rightsquigarrow$ Shows univariate marginal effect of feature j
- One plot per feature \rightsquigarrow Model is fully explainable via p additive plots



EBM WITH PAIRWISE INTERACTIONS

Generalized Additive Models plus Interactions (GA2M):

$$g(\mathbb{E}[y |]) = \theta_0 + \sum_{j=1}^p f_j(x_j) + \sum_{i < j} f_{ij}(x_i, x_j)$$





- **Motivation:** Univariate EBM does not model interactions
- **Challenge:** $O(p^2)$ potential pairwise interactions \rightsquigarrow often infeasible
- **Solution - FAST algorithm** ► Lou 2013:
 - Efficiently estimates importance of all feature pairs
 - Ranks pairs by reduction in residual sum of squares (RSS)
 - Avoids fitting EBM with each pairwise interaction
- **Result:**

Add only top-ranked interactions f_{ij} via a second-stage boosting step

\rightsquigarrow Performed after the univariate EBM has been trained
- **Interpretability preserved:** Each $f_{ij}(x_i, x_j)$ visualized as a 2D heatmap


FAST: PAIR-WISE INTERACTION STRENGTH

We evaluate a 4-leaf, axis-aligned tree T_{ij} over the 2D feature projection (x_i, x_j) .




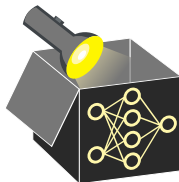
tree T_{ij} with 4 leaves

- 1 **Discretize** : Map each axis to $b \leq 256$ ordered bins (quantile or equal-width).




FAST: PAIR-WISE INTERACTION STRENGTH

We evaluate a 4-leaf, axis-aligned tree T_{ij} over the 2D feature projection (x_i, x_j) .



$\Rightarrow RSS_1$




$\Rightarrow RSS_2$

\vdots

\vdots

\vdots



$\Rightarrow RSS_{b^2}$

❶ **Discretize** : Map each axis to $b \leq 256$ ordered bins (quantile or equal-width).

❷ **Iterate** over b^2 candidate cuts (c_i, c_j) .

❸ **Fit** : For each cut, assign a constant $\hat{y}_r = \text{mean}(y \in r)$ to $r \in \{A, B, C, D\}$.

❹ **Compute RSS summed over all regions**:

$$RSS(c_i, c_j) = \sum_r \sum_{(x,y) \in r} (y - \hat{y}_r)^2$$
$$= \sum_r \left(\sum_{(x,y) \in r} y^2 - \frac{1}{n_r} \left(\sum_{(x,y) \in r} y \right)^2 \right)$$

❺ **Select** : Keep the split with minimal RSS.
 \rightsquigarrow largest RSS drop = strongest interaction.

FAST: USE RSS DROP

To evaluate a cut pair (c_i, c_j) , we use precomputed per-region statistics:

- For each region $r \in \{A, B, C, D\}$, compute:


$$S_r = \sum_{(x,y) \in r} y, \quad n_r = |\{(x,y) \in r\}|, \quad \hat{y}_r = S_r/n_r$$

- Plug into RSS summed over all regions:

$$\text{RSS}(c_i, c_j) = \sum_r \left(\sum_{(x,y) \in r} y^2 - \frac{1}{n_r} \left(\sum_{(x,y) \in r} y \right)^2 \right) = \sum_r \sum_{(x,y) \in r} y^2 + \sum_r \frac{S_r^2}{n_r}$$

- For a candidate cut, compute **RSS drop**:

$$\begin{aligned} \Delta \text{RSS}(c_i, c_j) &= \text{RSS}_{\text{parent}} - \text{RSS}(c_i, c_j) \\ &= \left(\sum_{i=1}^n (y^{(i)})^2 - \frac{S_n^2}{n} \right) - \sum_r \sum_{(x,y) \in r} y^2 + \sum_r \frac{S_r^2}{n_r} \end{aligned}$$



FAST: USE RSS DROP

Because $\sum_{i=1}^n (y^{(i)})^2 = \sum_r \sum_{(x,y) \in r} y^2$, *all squared target terms cancel*:


$$\Delta \text{RSS}(c_i, c_j) = \sum_r \frac{S_r^2}{n_r} - \frac{S_n^2}{n}$$

The parent term S_n^2/n is constant across all cuts. Hence


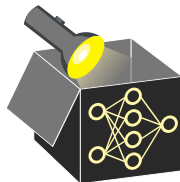

$$\textbf{maximize } \Delta \text{RSS}(c_i, c_j) = \sum_r \frac{S_r^2}{n_r} \iff \textbf{minimize } \text{RSS}(c_i, c_j).$$

Why is this efficient?

- Precompute cumulative sums of y and counts across the binned grid
- Enables fast lookup of region statistics S_r, n_r for any cut
- No additional data scan or recomputation needed across the b^2 candidate cuts
- For the best cut: Compare and select the largest $\Delta \text{RSS}(c_i, c_j)$.





EBM - BOOSTING PAIRWISE INTERACTIONS



- **Goal:** Fit each selected interaction $f_{ij}(x_i, x_j)$ on residuals from main effects
- Use tree-like predictor, inspired by FAST
 - Use two axis-aligned cuts (c_i, c_j)
 - Plus one refinement cut to increase flexibility while keeping interpretability
- Reuse region-wise sums from FAST lookup tables
- Greedy search for cut config minimizing RSS

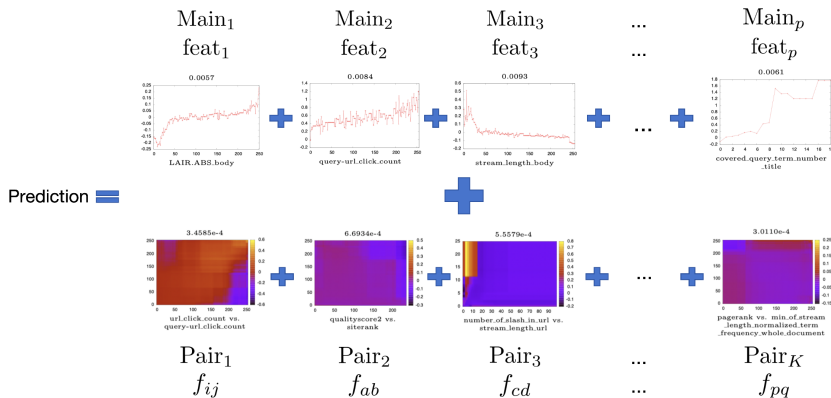

EBM - PREDICTION WITH PAIRWISE INTERACTIONS

- Each selected pair (x_i, x_j) is modeled by M boosted predictors trained on their residual interaction
- These are aggregated into a single bivariate function $f_{ij}(x_i, x_j)$
- The function is visualized as a 2D heatmap:
 - Axes: feature values of x_i and x_j
 - Color: contribution to the final prediction
 - Preserves human interpretability
- One heatmap is generated per selected pairwise interaction



EBM - FINAL MODEL STRUCTURE


- **Main effects:** One shape function $f_j(x_j)$ per feature (visualized as 1D plots)
- **Pairwise interactions:** Selected functions $f_{ij}(x_i, x_j)$ added for top K pairs (visualized as 2D heatmaps)
- **Prediction:** Additive sum of all univariate and selected bivariate contributions



EBM VS. MODEL-BASED BOOSTING

- Base learner

- **EBM**: bagged 2–4-leaf trees, *one feature* per tree \Rightarrow step-function shape f_j ▶ Lou 2012
- **MB-boost**: user chooses component-wise learner (linear term, P-spline, tree, random effect, ...) ▶ Hothorn 2007



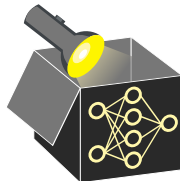
EBM VS. MODEL-BASED BOOSTING

- **Base learner**

- **EBM**: bagged 2–4-leaf trees, *one feature* per tree \Rightarrow step-function shape f_j ▶ Lou 2012
- **MB-boost**: user chooses component-wise learner (linear term, P-spline, tree, random effect, ...) ▶ Hothorn 2007

- **Iteration policy**

- **EBM**: round-robin ($\forall j$) each boosting pass; tiny learning rate $\eta \approx 0.01$.
- **MB-boost**: greedy; update the *single* component that yields the largest loss reduction.



EBM VS. MODEL-BASED BOOSTING

- **Base learner**


- **EBM**: bagged 2–4-leaf trees, *one feature* per tree \Rightarrow step-function shape f_j ▶ Lou 2012
- **MB-boost**: user chooses component-wise learner (linear term, P-spline, tree, random effect, ...) ▶ Hothorn 2007

- **Iteration policy**

- **EBM**: round-robin ($\forall j$) each boosting pass; tiny learning rate $\eta \approx 0.01$.
- **MB-boost**: greedy; update the *single* component that yields the largest loss reduction.

- **Regularisation**


- **EBM**: many iterations M (5–10k); early stopping via *internal* CV on out-of-bag samples; bagging further lowers variance.
- **MB-boost**: shrinkage $\nu \in (0, 1]$; early stop by CV/AIC; component selection acts like an L_0/L_1 penalty \rightarrow sparsity.



EBM VS. MODEL-BASED BOOSTING

- Interactions

- **EBM**: FAST ranks and selects top- K interaction pairs, fitted as bivariate trees \Rightarrow GA2M [▶ Lou 2013](#)
- **MB-boost**: interactions are modelled only when the user supplies dedicated interaction base learners; no automatic pairwise search



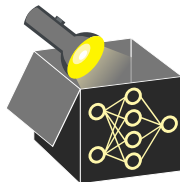
EBM VS. MODEL-BASED BOOSTING

- Interactions

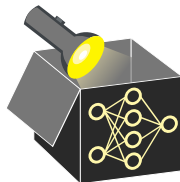
- **EBM**: FAST ranks and selects top- K interaction pairs, fitted as bivariate trees \Rightarrow GA2M ▶ Lou 2013
- **MB-boost**: interactions are modelled only when the user supplies dedicated interaction base learners; no automatic pairwise search

- Interpretability

- **EBM**:
 - one 1-D step plot for each f_j
 - small number of 2-D heat-maps for selected f_{ij}
- **MB-boost**: depends on selected learner: linear coefficients, smooth splines, random-effect curves, etc.



EBM VS. MODEL-BASED BOOSTING



- Interactions

- **EBM**: FAST ranks and selects top- K interaction pairs, fitted as bivariate trees \Rightarrow GA2M [▶ Lou 2013](#)
- **MB-boost**: interactions are modelled only when the user supplies dedicated interaction base learners; no automatic pairwise search

- Interpretability

- **EBM**:

- one 1-D step plot for each f_j
- small number of 2-D heat-maps for selected f_{ij}

- **MB-boost**: depends on selected learner: linear coefficients, smooth splines, random-effect curves, etc.

- Take-away

- *EBM* provides fast, interpretable, and interaction-sparse models
- *MB-boost* offers flexible stat modeling with built-in variable selection